

```
####In
## Name : Dhananjay Bokare
## Roll no:I4146
## Subject:LP-IV(DL)

In [1]:
import numpy as np
import pandas as pd
import random
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Conv2D, Dense, MaxPooling2D
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.datasets import mnist

In [2]:
(X_train, y_train), (X_test, y_test) = mnist.load_data()

In [3]:
print(X_train.shape)

(60000, 28, 28)

In [4]:
X_train[0].min(), X_train[0].max()

Out[4]: (0, 255)

In [5]:
X_train = (X_train - 0.0) / (255.0 - 0.0)
X_test = (X_test - 0.0) / (255.0 - 0.0)
X_train[0].min(), X_train[0].max()

Out[5]: (0.0, 1.0)

In [6]:
def plot_digit(image, digit, plt, i):
    plt.subplot(4, 5, i + 1)
    plt.imshow(image, cmap=plt.get_cmap('gray'))
    plt.title(f"Digit: {digit}")
    plt.xticks([])
    plt.yticks([])
plt.figure(figsize=(16, 10))
for i in range(20):
    plot_digit(X_train[i], y_train[i], plt, i)
plt.show()

Digit: 5
Digit: 0
Digit: 4
Digit: 1
Digit: 9
Digit: 2
Digit: 1
Digit: 3
Digit: 1
Digit: 4
Digit: 3
Digit: 5
Digit: 3
Digit: 6
Digit: 1
Digit: 7
Digit: 2
Digit: 8
Digit: 6
Digit: 9

In [7]:
X_train = X_train.reshape((X_train.shape + (1,)))
X_test = X_test.reshape((X_test.shape + (1,)))

In [8]:
y_train[0:20]

Out[8]: array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9],
      dtype=uint8)

In [9]:
model = Sequential([
    Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(100, activation="relu"),
    Dense(10, activation="softmax")
])

In [10]:
optimizer = SGD(learning_rate=0.01, momentum=0.9)
model.compile(
    optimizer=optimizer,
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
model.summary()

Model: "sequential"
Layer (type) Output Shape Param #
-----
conv2d (Conv2D) (None, 26, 26, 32) 320
max_pooling2d (MaxPooling2D) (None, 13, 13, 32) 0
flatten (Flatten) (None, 5408) 0
dense (Dense) (None, 100) 540900
dense_1 (Dense) (None, 10) 1010
-----
Total params: 542,230
Trainable params: 542,230
Non-trainable params: 0

In [11]:
model.fit(X_train, y_train, epochs=10, batch_size=32)

Epoch 1/10
1875/1875 [=====] - 19s 10ms/step - loss: 0.2392 - accuracy: 0.9275
Epoch 2/10
1875/1875 [=====] - 18s 9ms/step - loss: 0.0801 - accuracy: 0.9757
Epoch 3/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0511 - accuracy: 0.9846
Epoch 4/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0352 - accuracy: 0.9893
Epoch 5/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0268 - accuracy: 0.9917
Epoch 6/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0203 - accuracy: 0.9938
Epoch 7/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0145 - accuracy: 0.9955
Epoch 8/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0107 - accuracy: 0.9969
Epoch 9/10
1875/1875 [=====] - 19s 10ms/step - loss: 0.0092 - accuracy: 0.9975
Epoch 10/10
1875/1875 [=====] - 18s 10ms/step - loss: 0.0066 - accuracy: 0.9983

Out[11]: <keras.callbacks.History at 0x27b921bd40>

In [12]:
plt.figure(figsize=(16, 10))
for i in range(20):
    image = random.choice(X_test).squeeze()
    digit = np.argmax(model.predict(image.reshape((1, 28, 28, 1))))[0], axis=-1)
    plot_digit(image, digit, plt, i)
plt.show()

1/1 [=====] - 0s 105ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 20ms/step

Digit: 8
Digit: 9
Digit: 4
Digit: 8
Digit: 7
Digit: 2
Digit: 0
Digit: 3
Digit: 8
Digit: 8
Digit: 7
Digit: 3
Digit: 7
Digit: 9
Digit: 6
Digit: 7
Digit: 3
Digit: 2
Digit: 3
Digit: 0

In [13]:
predictions = np.argmax(model.predict(X_test), axis=-1)
accuracy_score(y_test, predictions)

313/313 [=====] - 1s 3ms/step

Out[13]: 0.9881

In [14]:
n=random.randint(0,9999)
plt.imshow(X_test[n])
plt.show()

0
5
10
15
20
25
0 5 10 15 20 25

In [15]:
predicted_value=model.predict(X_test)
print("Handwritten number in the image is= %d" %np.argmax(predicted_value[n]))

313/313 [=====] - 1s 3ms/step
Handwritten number in the image is= 0

In [16]:
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0]) #Test loss: 0.0296396646054
print('Test accuracy:', score[1])

Test loss: 0.03916610777378082
Test accuracy: 0.98809992275238

In [17]:
#The implemented CNN model is giving Loss=0.0462430156763683 and
#accuracy: 0.9872000217437744 for test mnist dataset
```